

# Mixing Logic and Deep Learning: The 'Logic as Loss Function' Approach

**Kyle Richardson**

Allen Institute for AI

June 2022

# The divide: logic vs. deep learning

**Classical logic:** designed for modeling *closed systems*.

# The divide: logic vs. deep learning

**Classical logic:** designed for modeling *closed systems*.

$$\begin{aligned}\forall x (0 \neq \mathbf{Succ}(x)) \\ \forall x, y (\mathbf{Succ}(x) = \mathbf{Succ}(y)) \rightarrow x = y \\ \forall x (x + 0 = x) \\ \forall x, y (x + \mathbf{Succ}(y) = \mathbf{Succ}(x + y)) \\ \forall x (x \cdot 0 = 0) \\ \forall \forall x, y (x \cdot \mathbf{Succ}(y) = x \cdot y + x)\end{aligned}$$

**Theorem:**  $\forall x, y (x \cdot y = \mathbf{Succ}(0) \rightarrow x = \mathbf{Succ}(0) \wedge y = \mathbf{Succ}(0))$

# The divide: logic vs. deep learning

**Classical logic:** designed for modeling *closed systems*.

$$\begin{aligned} &\forall x (0 \neq \mathbf{Succ}(x)) \\ &\forall x, y (\mathbf{Succ}(x) = \mathbf{Succ}(y)) \rightarrow x = y \\ &\forall x (x + 0 = x) \\ &\forall x, y (x + \mathbf{Succ}(y) = \mathbf{Succ}(x + y)) \\ &\forall x (x \cdot 0 = 0) \\ &\forall \forall x, y (x \cdot \mathbf{Succ}(y) = x \cdot y + x) \end{aligned}$$

**Theorem:**  $\forall x, y (x \cdot y = \mathbf{Succ}(0) \rightarrow x = \mathbf{Succ}(0) \wedge y = \mathbf{Succ}(0))$

**deep learning:** Used for modeling open systems, situations involving *partial information*, making (fast) predictions vs. explicit reasoning.

# The divide: logic vs. deep learning

**Classical logic:** designed for modeling *closed systems*.

$$\begin{aligned}\forall x(0 \neq \mathbf{Succ}(x)) \\ \forall x, y(\mathbf{Succ}(x) = \mathbf{Succ}(y)) \rightarrow x = y \\ \forall x(x + 0 = x) \\ \forall x, y(x + \mathbf{Succ}(y) = \mathbf{Succ}(x + y)) \\ \forall x(x \cdot 0 = 0) \\ \forall x, y(x \cdot \mathbf{Succ}(y) = x \cdot y + x)\end{aligned}$$

**Theorem:**  $\forall x, y(x \cdot y = \mathbf{Succ}(0) \rightarrow x = \mathbf{Succ}(0) \wedge y = \mathbf{Succ}(0))$

**deep learning:** Used for modeling open systems, situations involving *partial information*, making (fast) predictions vs. explicit reasoning.

**Premise:** A man with a hat is riding his bicycle down the street.

**Hypothesis:** A person is moving with the help of their legs.

**Prediction:** Entailment

**Premise:** A person is moving with the help of their legs.

**Hypothesis:** A person is moving

**Prediction:** Entailment

Two conceptual tools for relating logic and deep learning

# 1. Predictions-as-propositions

**x1:** A man with a hat is riding his bicycle down the street.

**x2:** A person is moving with the help of their legs.

**Prediction:** Entailment

**x2:** A person is moving with the help of their legs.

**x3:** A person is moving

**Prediction:** Entailment

# 1. Predictions-as-propositions

**x1:** A man with a hat is riding his bicycle down the street.

**x2:** A person is moving with the help of their legs.

**Prediction:** **Entailment**

**x2:** A person is moving with the help of their legs.

**x3:** A person is moving

**Prediction:** **Entailment**





# 1. Predictions-as-propositions

**x1:** A man with a hat is riding his bicycle down the street.

**x2:** A person is moving with the help of their legs.

**Prediction:** **Entailment**

**x2:** A person is moving with the help of their legs.

**x3:** A person is moving

**Prediction:** **Entailment**



# 1. Predictions-as-propositions

**x1:** A man with a hat is riding his bicycle down the street.

**x2:** A person is moving with the help of their legs.


**Prediction:** **Entailment**

**x2:** A person is moving with the help of their legs.


**x3:** A person is moving

**Prediction:** **Entailment**

**symbolically:**  $E(x_1, x_2)$



**symbolically:**  $E(x_2, x_3)$



# 1. Predictions-as-propositions

**x1:** A man with a hat is riding his bicycle down the street.

**x2:** A person is moving with the help of their legs.


**Prediction:** **Entailment**

**x2:** A person is moving with the help of their legs.


**x3:** A person is moving

**Prediction:** **Entailment**

**symbolically:**  $E(x_1, x_2)$




**symbolically:**  $E(x_2, x_3)$



Proposition	Meaning
$E(x, y)$	$x$ entails $y$
$C(x, y)$	$x$ contradicts $y$


[Notation from [Li et al. \(2019\)](#)]

# 1. Predictions-as-propositions




Why is this helpful?

# 1. Predictions-as-propositions




Why is this helpful?




# Distinguishing good, bad and very bad model predictions

constraints (*our expectations*)

Minervini and Riedel (2018); Li et al. (2019)




# Distinguishing good, bad and very bad model predictions



The different predictions a model (or set of models) might make:

# Distinguishing good, bad and very bad model predictions




The different predictions a model (or set of models) might make:

$x_1$  : *A man with a hat is riding his bicycle down the street* entails  $x_2$  : *A person is moving with the help of their legs*,  $x_3$  : entails *A person is moving*

predictions	predictions correct	predictions consistent
$\mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{E}(x_1, x_3)$	✓ (3/3)	✓ (3/3)
$\mathbf{C}(x_1, x_2) \wedge \mathbf{C}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (0/3)	✓ (3/3)
$\mathbf{E}(x_1, x_2) \wedge \mathbf{C}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (1/3)	✓ (3/3)
$\mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (2/3)	✗ (0/3)



# Distinguishing good, bad and very bad model predictions




The different predictions a model (or set of models) might make:

$x_1$  : *A man with a hat is riding his bicycle down the street* entails  $x_2$  : *A person is moving with the help of their legs*,  $x_3$  : entails *A person is moving*

predictions	predictions correct	predictions consistent
$E(x_1, x_2) \wedge E(x_2, x_3) \wedge E(x_1, x_3)$	✓ (3/3)	✓ (3/3)
$C(x_1, x_2) \wedge C(x_2, x_3) \wedge C(x_1, x_3)$	✗ (0/3)	✓ (3/3)
$E(x_1, x_2) \wedge C(x_2, x_3) \wedge C(x_1, x_3)$	✗ (1/3)	✓ (3/3)
$E(x_1, x_2) \wedge E(x_2, x_3) \wedge C(x_1, x_3)$	✗ (2/3)	✗ (0/3)

Thinking of predictions as **symbolic objects**; brings rigor to interpreting model behavior, can clarify what we mean by ‘good’ vs. ‘bad’.

# Distinguishing good, bad and very bad model predictions



The different predictions a model (or set of models) might make:

$x_1$  : *A man with a hat is riding his bicycle down the street* entails  $x_2$  : *A person is moving with the help of their legs*,  $x_3$  : entails *A person is moving*

predictions	predictions correct	
$E(x_1, x_2) \wedge E(x_2, x_3) \wedge E(x_1, x_3)$	✓ (3/3)	
$C(x_1, x_2) \wedge C(x_2, x_3) \wedge C(x_1, x_3)$	✗ (0/3)	
$E(x_1, x_2) \wedge C(x_2, x_3) \wedge C(x_1, x_3)$	✗ (1/3)	
$E(x_1, x_2) \wedge E(x_2, x_3) \wedge C(x_1, x_3)$	✗ (2/3)	

Thinking of predictions as **symbolic objects**; brings rigor to interpreting model behavior, can clarify what we mean by 'good' vs. 'bad'.

# Annotations-as-logical-specifications: what we expect

**Question:**  $(q)$  Where is a frisbee in play likely to be?  $(m)$  1) **air** 2) ...

**Prediction**  $(q, m \rightarrow a)$ :  $(a)$  "air"

**Prediction**  $(q, m \rightarrow e, a)$ :  $(e)$  "A frisbee floats on air", "air"

**Prediction**  $(q, m \rightarrow p)$ :  $(p)$  "A frisby in play is likely to be in the air"

# Annotations-as-logical-specifications: what we expect

**Question:**  $(q)$  Where is a frisbee in play likely to be?  $(m)$  1) **air** 2) ...

**Prediction**  $(q, m \rightarrow a)$ :  $(a)$  "air"

**Prediction**  $(q, m \rightarrow e, a)$ :  $(e)$  "A frisbee floats on air", "air"

**Prediction**  $(q, m \rightarrow p)$ :  $(p)$  "A frisby in play is likely to be in the air"

(Labeled) datasets specify what we want a model to do and learn.

instance	meaning	proposition (logic)
$(q, m, a)$	$a$ is the correct answer to $q$ in $m$	$Q(q, a)$
$(q, m, e, a)$	$e$ is the correct <u>explanation</u> of $q$ with answer $a$	$Ex(q, e + a)$
$(q, m, p)$	$p$ is the correct <u>proposition</u> corresponding to $q$ and $a$ .	$P(q, p)$

# Annotations-as-logical-specifications: what we expect

**Question:** (q) Where is a frisbee in play likely to be? (m) 1) air 2) ...

**Prediction** (q, m  $\rightarrow$  a): (a) "air"

**Prediction** (q, m  $\rightarrow$  e, a): (e) "A frisbee floats on air", "air"

**Prediction** (q, m  $\rightarrow$  p): (p) "A frisby in play is likely to be in the air"

(Labeled) datasets specify what we want a model to do and learn.

instance	meaning	proposition (logic)
(q, m, a)	a is the correct answer to q in m	$Q(q, a)$
(q, m, e, a)	e is the correct <u>explanation</u> of q with answer a	$Ex(q, e + a)$
(q, m, p)	p is the correct <u>proposition</u> corresponding to q and a.	$P(q, p)$

We can also think of annotations as logical propositions and formulae.

$$Q(q, a) \wedge Ex(q, e + a) \wedge P(q, p)$$

# Annotations-as-logic-specifications: what we expect

**Question:**  $(q)$  Where is a frisbee in play likely to be?  $(m)$  1) air 2) ...

**Prediction**  $(q, m \rightarrow a)$ :  $(a)$  "air"

**Prediction**  $(q, m \rightarrow e, a)$ :  $(e)$  "A frisbee floats on air", "air"

**Prediction**  $(q, m \rightarrow p)$ :  $(p)$  "A frisby in play is likely to be in the air"

$$\underbrace{Q(q, a) \wedge \text{Ex}(q, e + a) \wedge P(q, p)}_{\text{should make correct predictions}} \wedge \underbrace{\forall q, p (P(q, p) \rightarrow \text{Bel}(p))}_{\text{should believe propositions}} \wedge \underbrace{\forall q, a, e (\text{Ex}(q, e + a) \rightarrow Q(q + e, a))}_{\text{Answer should be invariant given its explanation}}$$

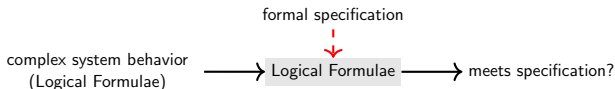
**Why is this helpful?** Be more clear about what we expect, understand gap between what we expect and what we *actually* do, imagine new tasks.

## Conceptual Tools: Predictions-as-propositions, Annotations-as-specifications

Thinking of *model predictions as symbolic objects*; *annotations and our expectations* as logical formulas.

# Conceptual Tools: Predictions-as-propositions, Annotations-as-specifications

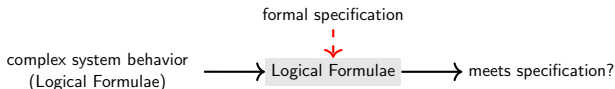
Thinking of *model predictions as symbolic objects*; *annotations and our expectations* as logical formulas.





# Conceptual Tools: Predictions-as-propositions, Annotations-as-specifications

Thinking of *model predictions as symbolic objects*; *annotations and our expectations* as logical formulas.



**note:** These are just conceptual tools, not particularly useful yet.

Technical Tool: Multi-valued Logic and the ‘Logic as Loss Function’ approach (see review in [Marra et al. \(2021\)](#)).

# Multi-valued logic

Logical propositions come in different flavors

<b>Boolean Propositions</b>	$\mathbf{P} \in \{0, 1\}$
<b>(Probabilistic) Boolean Propositions</b>	$\mathbf{P}$ is 1 with prob. $p \in [0, 1]$ (0 with prob. $1 - p$ )
	non-classical logic
<b>(Finite-)N-Valued Propositions</b>	$\mathbf{P} \in \{0, 1, \dots, N\}$
<b>Real-Valued (Fuzzy) Propositions</b>	$\mathbf{P} \in [0, 1]$ ( <i>truth degree</i> ),

# Multi-valued logic

Logical propositions come in different flavors

<b>Boolean Propositions</b>	$P \in \{0, 1\}$
<b>(Probabilistic) Boolean Propositions</b>	$P$ is 1 with prob. $p \in [0, 1]$ (0 with prob. $1 - p$ )
	non-classical logic
<b>(Finite-)N-Valued Propositions</b>	$P \in \{0, 1, \dots, N\}$
<b>Real-Valued (Fuzzy) Propositions</b>	$P \in [0, 1]$ ( <i>truth degree</i> ),

**Fuzzy Logic:** Logical operators ( $\wedge, \vee, \neg, \rightarrow$ ) are turned into real-valued functions (t-norms); generalizes Boolean logic to continuous values.

# Multi-valued logic

Logical propositions come in different flavors

<b>Boolean Propositions</b>	$P \in \{0, 1\}$
<b>(Probabilistic) Boolean Propositions</b>	$P$ is 1 with prob. $p \in [0, 1]$ (0 with prob. $1 - p$ )
	non-classical logic
<b>(Finite-)N-Valued Propositions</b>	$P \in \{0, 1, \dots, N\}$
<b>Real-Valued (Fuzzy) Propositions</b>	$P \in [0, 1]$ ( <i>truth degree</i> ),

**Fuzzy Logic:** Logical operators ( $\wedge, \vee, \neg, \rightarrow$ ) are turned into real-valued functions (t-norms); generalizes Boolean logic to continuous values.

see [Li et al. \(2019\)](#); [Grespan et al. \(2021\)](#)

	Boolean Logic	Product	Łukasiewicz	Gödel
T-norm	$P_1 \wedge P_2$	$P_1 \cdot P_2$	$\max(0, P_1 + P_2 - 1)$	$\min(P_1, P_2)$
T-conorm	$P_1 \vee P_2$	$P_1 + P_2 - P_1 \cdot P_2$	$\min(1, P_1 + P_2)$	$\max(P_1, P_2)$
Negation	$\neg P$	$1 - P$	$1 - P$	$1 - P$
Residuum	$P_1 \rightarrow P_2$	$\min(1, \frac{P_2}{P_1})$	$\min(1, 1 - P_1 + P_2)$	$P_2$

# Multi-valued logic

Logical propositions come in different flavors

<b>Boolean Propositions</b>	$P \in \{0, 1\}$
<b>(Probabilistic) Boolean Propositions</b>	$P$ is 1 with prob. $p \in [0, 1]$ (0 with prob. $1 - p$ )
	non-classical logic
<b>(Finite-)N-Valued Propositions</b>	$P \in \{0, 1, \dots, N\}$
<b>Real-Valued (Fuzzy) Propositions</b>	$P \in [0, 1]$ ( <i>truth degree</i> ),

**Fuzzy Logic:** Logical operators ( $\wedge, \vee, \neg, \rightarrow$ ) are turned into real-valued functions (t-norms); generalizes Boolean logic to continuous values.

see [Li et al. \(2019\)](#); [Grespan et al. \(2021\)](#)

	Boolean Logic	Product	Łukasiewicz	Gödel
T-norm	$P_1 \wedge P_2$	$P_1 \cdot P_2$	$\max(0, P_1 + P_2 - 1)$	$\min(P_1, P_2)$
T-conorm	$P_1 \vee P_2$	$P_1 + P_2 - P_1 \cdot P_2$	$\min(1, P_1 + P_2)$	$\max(P_1, P_2)$
Negation	$\neg P$	$1 - P$	$1 - P$	$1 - P$
Residuum	$P_1 \rightarrow P_2$	$\min(1, \frac{P_2}{P_1})$	$\min(1, 1 - P_1 + P_2)$	$P_2$

**Note:** *At the extremes 0 and 1, work exactly as classical counterparts.*

# Turning Specifications into Loss Functions


	Boolean Logic	Product	Łukasiewicz	Gödel
T-norm	$P_1 \wedge P_2$	$P_1 \cdot P_2$	$\max(0, P_1 + P_2 - 1)$	$\min(P_1, P_2)$
T-conorm	$P_1 \vee P_2$	$P_1 + P_2 - P_1 \cdot P_2$	$\min(1, P_1 + P_2)$	$\max(P_1, P_2)$
Negation	$\neg P$	$1 - P$	$1 - P$	$1 - P$
Residuum	$P_1 \rightarrow P_2$	$\min(1, \frac{P_2}{P_1})$	$\min(1, 1 - P_1 + P_2)$	$P_2$

**Why are these (T-norms) useful?**

# Turning Specifications into Loss Functions

	Boolean Logic	Product	Łukasiewicz	Gödel
T-norm	$P_1 \wedge P_2$	$P_1 \cdot P_2$	$\max(0, P_1 + P_2 - 1)$	$\min(P_1, P_2)$
T-conorm	$P_1 \vee P_2$	$P_1 + P_2 - P_1 \cdot P_2$	$\min(1, P_1 + P_2)$	$\max(P_1, P_2)$
Negation	$\neg P$	$1 - P$	$1 - P$	$1 - P$
Residuum	$P_1 \rightarrow P_2$	$\min(1, \frac{P_2}{P_1})$	$\min(1, 1 - P_1 + P_2)$	$P_2$

Why are these (T-norms) useful?





# Turning Specifications into Loss Functions

	Boolean Logic	Product	Łukasiewicz	Gödel
T-norm	$P_1 \wedge P_2$	$P_1 \cdot P_2$	$\max(0, P_1 + P_2 - 1)$	$\min(P_1, P_2)$
T-conorm	$P_1 \vee P_2$	$P_1 + P_2 - P_1 \cdot P_2$	$\min(1, P_1 + P_2)$	$\max(P_1, P_2)$
Negation	$\neg P$	$1 - P$	$1 - P$	$1 - P$
Residuum	$P_1 \rightarrow P_2$	$\min(1, \frac{P_2}{P_1})$	$\min(1, 1 - P_1 + P_2)$	$P_2$

Why are these (T-norms) useful?

Specifications		
$Q(q, a) \wedge \text{Ex}(q, e) \wedge P(q, p) \wedge$	$(P(q, p) \rightarrow \text{Bel}(p)) \wedge$	$(\text{Ex}(q, e + a) \rightarrow Q(q + e, a))$
should make correct predictions	should believe propositions	Answer should be invariant given its explanation

(translation to T-norms)

Differentiable Loss Function

$$\text{Loss}_{\text{model}} \left( Q(q, a) \wedge \text{Ex}(q, e) \wedge \dots \right) = \underbrace{\sum_{P \in \{Q(q, a), \text{Ex}(q, e), \dots\}} -\log \underbrace{p_{\text{model}}(P)}_{\text{fuzzy truth degree}}}_{\text{Product T-norm } (-\log \text{ prob.})}$$

# Training Objectives as Logical Specifications

$$\underbrace{\mathbf{Q}(q, a) \wedge \mathbf{Ex}(q, e) \wedge \mathbf{P}(q, p)}_{\text{Atomic predictions}} \wedge \underbrace{(\mathbf{P}(q, p) \rightarrow \mathbf{Bel}(p))}_{\text{should believe propositions}} \wedge \underbrace{(\mathbf{Ex}(q, e + a) \rightarrow \mathbf{Q}(q + e, a))}_{\text{Answer should be invariant given its explanation}}$$

Additional Constraints

We can think of a (supervised) dataset  $D = \{(x_j, y_j)\}_{j=0}^N$  as a set of true atomic propositions  $D_p = \{\mathbf{Y}_1, \dots, \mathbf{Y}_N\}$  with constraints  $C$ .

Goal	Logical Formula	Loss Function (Product)
Make Correct Predictions	$\bigwedge_{\mathbf{Y} \in D_p} \mathbf{Y}$	$\sum_{\mathbf{Y} \in D_p} -\log p_{\text{model}}(\mathbf{Y})$
Believe your propositions	$\bigwedge_{\mathbf{P}(q, p) \in D_p} \mathbf{P}(q, p) \rightarrow \mathbf{Bel}(p)$	$\sum_{\mathbf{P}(q, p) \in D_p} \text{ReLU}(\log p_{\text{model}}(\mathbf{P}(q, p)) - \log p_{\text{model}}(\mathbf{Bel}(p)))$

# Training Objectives as Logical Specifications

$$\boxed{
 \underbrace{Q(q, a) \wedge \text{Ex}(q, e) \wedge P(q, p)}_{\text{Atomic predictions}} \wedge \underbrace{(P(q, p) \rightarrow \text{Bel}(p))}_{\text{should believe propositions}} \wedge \underbrace{(\text{Ex}(q, e + a) \rightarrow Q(q + e, a))}_{\text{Answer should be invariant given its explanation}}$$

Additional Constraints

We can think of a (supervised) dataset  $D = \{(x_j, y_j)\}_{j=0}^N$  as a set of true atomic propositions  $D_p = \{\mathbf{Y}_1, \dots, \mathbf{Y}_N\}$  with constraints  $C$ .

Goal	Logical Formula	Loss Function (Product)
<i>Make Correct Predictions</i>	$\bigwedge_{\mathbf{Y} \in D_p} \mathbf{Y}$	$\sum_{\mathbf{Y} \in D_p} -\log p_{\text{model}}(\mathbf{Y})$
<i>Believe your propositions</i>	$\bigwedge_{P(q,p) \in D_p} P(q, p) \rightarrow \text{Bel}(p)$	$\sum_{P(q,p) \in D_p} \text{ReLU}(\log p_{\text{model}}(P(q, p)) - \log p_{\text{model}}(\text{Bel}(p)))$

**Observation** (Rocktäschel et al., 2015; Li et al., 2019): Translating product conjunction to negative log space yields ordinary **cross-entropy** loss.

# Training Objectives as Logical Specifications

$$\boxed{
 \underbrace{Q(q, a) \wedge \text{Ex}(q, e) \wedge P(q, p)}_{\text{Atomic predictions}} \wedge \underbrace{(P(q, p) \rightarrow \text{Bel}(p))}_{\text{should believe propositions}} \wedge \underbrace{(\text{Ex}(q, e + a) \rightarrow Q(q + e, a))}_{\text{Answer should be invariant given its explanation}}$$

Additional Constraints

We can think of a (supervised) dataset  $D = \{(x_j, y_j)\}_{j=0}^N$  as a set of true atomic propositions  $D_p = \{\mathbf{Y}_1, \dots, \mathbf{Y}_N\}$  with constraints  $C$ .

Goal	Logical Formula	Loss Function (Product)
<i>Make Correct Predictions</i>	$\bigwedge_{\mathbf{Y} \in D_p} \mathbf{Y}$	$\sum_{\mathbf{Y} \in D_p} -\log p_{\text{model}}(\mathbf{Y})$
<i>Believe your propositions</i>	$\bigwedge_{P(q,p) \in D_p} P(q, p) \rightarrow \text{Bel}(p)$	$\sum_{P(q,p) \in D_p} \text{ReLU}(\log p_{\text{model}}(P(q, p)) - \log p_{\text{model}}(\text{Bel}(p)))$

**Observation:** There is often a large gap between what we train models to do (e.g., **make correct predictions**) and we expect models to know.

# Logic as Loss Function, Logic in the Weights


Logical constraints serve as regularizer, *soft constraints* over hypothesis space that favor solutions closer to knowledge (undirected models).

$$\mathbf{Loss} = \underbrace{\mathbf{Loss}_{\text{CE}}}_{\text{ordinary loss}} + \lambda \underbrace{\mathbf{Loss}_{\text{constraints}}}_{\text{logical constraints}}$$

# Logic as Loss Function, Logic in the Weights

Logical constraints serve as regularizer, *soft constraints* over hypothesis space that favor solutions closer to knowledge (undirected models).

$$\text{Loss} = \underbrace{\text{Loss}_{\text{CE}}}_{\text{ordinary loss}} + \underbrace{\lambda \text{Loss}_{\text{constraints}}}_{\text{logical constraints}}$$




predictions	predictions correct	predictions consistent	prediction loss	constraint loss
$\mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{E}(x_1, x_3)$	✓ (3/3)	✓ (3/3)	low	low
$\mathbf{C}(x_1, x_2) \wedge \mathbf{C}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (0/3)	✓ (3/3)	high	low
$\mathbf{E}(x_1, x_2) \wedge \mathbf{C}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (1/3)	✓ (3/3)	high	low
$\mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (2/3)	✗ (0/3)	medium	high

# Logic as Loss Function, Logic in the Weights

Logical constraints serve as regularizer, *soft constraints* over hypothesis space that favor solutions closer to knowledge (undirected models).

$$\text{Loss} = \underbrace{\text{Loss}_{\text{CE}}}_{\text{ordinary loss}} + \underbrace{\lambda \text{Loss}_{\text{constraints}}}_{\text{logical constraints}}$$




predictions	predictions correct	predictions consistent	prediction loss	constraint loss
$\mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{E}(x_1, x_3)$	✓ (3/3)	✓ (3/3)	low	low
$\mathbf{C}(x_1, x_2) \wedge \mathbf{C}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (0/3)	✓ (3/3)	high	low
$\mathbf{E}(x_1, x_2) \wedge \mathbf{C}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (1/3)	✓ (3/3)	high	low
$\mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$	✗ (2/3)	✗ (0/3)	medium	high

**Observation:** constraint loss doesn't contribute much outside of penalizing the last case; need to be carefully constructed.

# Logic as Loss Function, Logic in the Weights

Logical constraints serve as regularizer, *soft constraints* over hypothesis space that favor solutions closer to knowledge (undirected models).

$$\text{Loss} = \underbrace{\text{Loss}_{\text{CE}}}_{\text{ordinary loss}} + \underbrace{\lambda \text{Loss}_{\text{constraints}}}_{\text{logical constraints}}$$







predictions		predictions consistent		constraint loss
$\text{E}(x_1, x_2) \wedge \text{E}(x_2, x_3) \wedge \text{E}(x_1, x_3)$		✓ (3/3)		low
$\text{C}(x_1, x_2) \wedge \text{C}(x_2, x_3) \wedge \text{C}(x_1, x_3)$		✓ (3/3)		low
$\text{E}(x_1, x_2) \wedge \text{C}(x_2, x_3) \wedge \text{C}(x_1, x_3)$		✓ (3/3)		low
$\text{E}(x_1, x_2) \wedge \text{E}(x_2, x_3) \wedge \text{C}(x_1, x_3)$		✗ (0/3)		high

**Practical concerns:** ensure consistency loss doesn't overwhelm your prediction loss; many tricks for this (loss weighting  $\lambda$ , annealing).



# Does this Help?

- Li et al. (2019) apply to NLI, focus on basic order relations between inference types (transitivity, symmetry), study different t-norm approaches.



# Does this Help?

- ▶ **NLP:** NLI ([Minervini and Riedel, 2018](#); [Li et al., 2019](#)), question-answering ([Asai and Hajishirzi, 2020](#)), relation extraction ([Rocktäschel et al., 2015](#)), other ([Grespan et al., 2021](#)).

# Does this Help?

- ▶ **NLP:** NLI ([Minervini and Riedel, 2018](#); [Li et al., 2019](#)), question-answering ([Asai and Hajishirzi, 2020](#)), relation extraction ([Rocktäschel et al., 2015](#)), other ([Grespan et al., 2021](#)).

Can significantly improve consistency and training efficiency, mixed results on improving end-task performance, though.

# Does this Help?

- ▶ **NLP: NLI** ([Minervini and Riedel, 2018](#); [Li et al., 2019](#)), question-answering ([Asai and Hajishirzi, 2020](#)), relation extraction ([Rocktäschel et al., 2015](#)), other ([Grespan et al., 2021](#)).

Can significantly improve consistency and training efficiency, mixed results on improving end-task performance, though.

- ▶ Widely used elsewhere in neural-symbolic modeling (see [Marra et al. \(2021\)](#)), many open technical issues (see [Grespan et al. \(2021\)](#)).

## An Alternative Tool: Weighted Model Counting

## An Alternative Tool: Weighted Model Counting

*(The proper way to do things!)*

# A Probabilistic Approach

- ▶ The semantics of Fuzzy logic has issues, not easy to translate back to ordinary Boolean logic, not amenable to probabilistic inference.

**Boolean Propositions**

$P \in \{0, 1\}$

**(Probabilistic) Boolean Propositions**

$P$  is 1 with prob.  $p \in [0, 1]$  (0 with prob.  $1 - p$ )

**Real-Valued (Fuzzy) Propositions**

$P \in [0, 1]$ , *truth degree*  $t(P)$

# A Probabilistic Approach

- ▶ The semantics of Fuzzy logic has issues, not easy to translate back to ordinary Boolean logic, not amenable to probabilistic inference.

Boolean Propositions	$P \in \{0, 1\}$
(Probabilistic) Boolean Propositions	$P$ is 1 with prob. $p \in [0, 1]$ (0 with prob. $1 - p$ )
Real-Valued (Fuzzy) Propositions	$P \in [0, 1]$ , truth degree $t(P)$

**Example:** Assume we have a proposition  $P$  with a weight 0.3 and we want to get a weight for  $P \wedge P$  (see (Marra et al., 2021, p43))



# A Probabilistic Approach

- ▶ The semantics of Fuzzy logic has issues, not easy to translate back to ordinary Boolean logic, not amenable to probabilistic inference.

Boolean Propositions	$P \in \{0, 1\}$
(Probabilistic) Boolean Propositions	$P$ is 1 with prob. $p \in [0, 1]$ (0 with prob. $1 - p$ )
Real-Valued (Fuzzy) Propositions	$P \in [0, 1]$ , truth degree $t(P)$

**Example:** Assume we have a proposition  $P$  with a weight 0.3 and we want to get a weight for  $P \wedge P$  (see (Marra et al., 2021, p43))

$$\begin{aligned} p(P \wedge P) &= p(P) && \text{(possible world semantics)} \\ t(P \wedge P) &= t(P) \times t(P) = 0.15 && \text{(product t-norm semantics)} \end{aligned}$$

## A Probabilistic Approach: Possible World Semantics

**Three propositions:**  $0.9 :: P_1$  (*A dog is a type of mammal*),  $0.8 :: P_2$  (*A mammal is a type of animal*),  $0.45 :: P_3$  (*A bulldog is a dog*)

world $W$	$P_1$	$P_2$	$P_3$	$p(W)$
$W_1$	0	0	0	$(0.1 \times 0.2 \times 0.55) = 0.01$
$W_2$	1	1	1	$(0.9 \times 0.8 \times 0.45) = 0.32$
$W_3$	0	0	1	$(0.1 \times 0.2 \times 0.45) = 0.009$
$W_4$	0	1	1	$(0.1 \times 0.8 \times 0.45) = 0.036$
$W_5$	1	0	1	$(0.9 \times 0.2 \times 0.45) = 0.081$
$W_6$	1	1	0	$(0.9 \times 0.8 \times 0.55) = 0.39$
$W_7$	0	1	0	$(0.1 \times 0.8 \times 0.55) = 0.04$
$W_8$	1	0	0	$(0.9 \times 0.2 \times 0.55) = 0.08$

# A Probabilistic Approach: Possible World Semantics

**Three propositions:**  $0.9 :: P_1$  (*A dog is a type of mammal*),  $0.8 :: P_2$  (*A mammal is a type of animal*),  $0.45 :: P_3$  (*A bulldog is a dog*)

world $W$	$P_1$	$P_2$	$P_3$	$p(W)$
$W_1$	0	0	0	$(0.1 \times 0.2 \times 0.55) = 0.01$
$W_2$	1	1	1	$(0.9 \times 0.8 \times 0.45) = 0.32$
$W_3$	0	0	1	$(0.1 \times 0.2 \times 0.45) = 0.009$
$W_4$	0	1	1	$(0.1 \times 0.8 \times 0.45) = 0.036$
$W_5$	1	0	1	$(0.9 \times 0.2 \times 0.45) = 0.081$
$W_6$	1	1	0	$(0.9 \times 0.8 \times 0.55) = 0.39$
$W_7$	0	1	0	$(0.1 \times 0.8 \times 0.55) = 0.04$
$W_8$	1	0	0	$(0.9 \times 0.2 \times 0.55) = 0.08$

$$p(W) = \underbrace{\prod_{P \in W^1} p(P)}_{\text{true in } W, W^1} \times \underbrace{\prod_{P \in W^0} 1 - p(P)}_{\text{false in } W, W^0}$$

$$p_{\text{query}}(P) = \sum_{W \text{ s.t. } W \models P} p(W)$$

# A Probabilistic Approach: Possible World Semantics

**Three propositions:**  $0.9 :: P_1$  (*A dog is a type of mammal*),  $0.8 :: P_2$  (*A mammal is a type of animal*),  $0.45 :: P_3$  (*A bulldog is a dog*)

world $W$	$P_1$	$P_2$	$P_3$	$p(W)$
$W_1$	0	0	0	$(0.1 \times 0.2 \times 0.55) = 0.01$
$W_2$	1	1	1	$(0.9 \times 0.8 \times 0.45) = 0.32$
$W_3$	0	0	1	$(0.1 \times 0.2 \times 0.45) = 0.009$
$W_4$	0	1	1	$(0.1 \times 0.8 \times 0.45) = 0.036$
$W_5$	1	0	1	$(0.9 \times 0.2 \times 0.45) = 0.081$
$W_6$	1	1	0	$(0.9 \times 0.8 \times 0.55) = 0.39$
$W_7$	0	1	0	$(0.1 \times 0.8 \times 0.55) = 0.04$
$W_8$	1	0	0	$(0.9 \times 0.2 \times 0.55) = 0.08$

$$p(W) = \underbrace{\prod_{P \in W^1} p(P)}_{\text{true in } W, W^1} \times \underbrace{\prod_{P \in W^0} 1 - p(P)}_{\text{false in } W, W^0}$$

$$p_{\text{query}}(P) = \sum_{W \text{ s.t. } W \models P} p(W)$$

# A Probabilistic Approach: Possible World Semantics

**Three propositions:**  $0.9 :: P_1$  (*A dog is a type of mammal*),  $0.8 :: P_2$  (*A mammal is a type of animal*),  $0.45 :: P_3$  (*A bulldog is a dog*)

world $W$	$P_1$	$P_2$	$P_3$	$p(W)$
$W_1$	0	0	0	$(0.1 \times 0.2 \times 0.55) = 0.01$
$W_2$	1	1	1	$(0.9 \times 0.8 \times 0.45) = 0.32$
$W_3$	0	0	1	$(0.1 \times 0.2 \times 0.45) = 0.009$
$W_4$	0	1	1	$(0.1 \times 0.8 \times 0.45) = 0.036$
$W_5$	1	0	1	$(0.9 \times 0.2 \times 0.45) = 0.081$
$W_6$	1	1	0	$(0.9 \times 0.8 \times 0.55) = 0.39$
$W_7$	0	1	0	$(0.1 \times 0.8 \times 0.55) = 0.04$
$W_8$	1	0	0	$(0.9 \times 0.2 \times 0.55) = 0.08$

$$p(W) = \underbrace{\prod_{P \in W^1} p(P)}_{\text{true in } W, W^1} \times \underbrace{\prod_{P \in W^0} 1 - p(P)}_{\text{false in } W, W^0}$$

$$p_{\text{query}}(P) = \sum_{W \text{ s.t. } W \models P} p(W)$$

# A Probabilistic Approach: Weighted Model Counting

- **Querying:** generalizes to any propositional formula  $\alpha$  and reducible to the problem of Weighted Model Counting (WMC) ([Chavira and Darwiche, 2008](#))

$$p_{query}(\alpha) = \underbrace{\sum_{W \text{ s.t. } W \models \alpha} p(W)}_{\text{WMC}}$$

# A Probabilistic Approach: Weighted Model Counting

- **Querying:** generalizes to any propositional formula  $\alpha$  and reducible to the problem of Weighted Model Counting (WMC) (Chavira and Darwiche, 2008)

$$p_{\text{query}}(\alpha) = \underbrace{\sum_{W \text{ s.t. } W \models \alpha} p(W)}_{\text{WMC}}$$

constraints

$$\mathbf{P}_1 \rightarrow \mathbf{P}_2 \wedge \mathbf{P}_2 \rightarrow \mathbf{P}_3$$

world $W$	$\mathbf{P}_1$	$\mathbf{P}_2$	$\mathbf{P}_3$	model?	$p(W)$
$W_1$	0	0	0	yes	$(0.1 \times 0.2 \times 0.55) = 0.01$
$W_2$	1	1	1	yes	$(0.9 \times 0.8 \times 0.45) = 0.32$
$W_3$	0	0	1	yes	$(0.1 \times 0.2 \times 0.45) = 0.009$
$W_4$	0	1	1	yes	$(0.1 \times 0.8 \times 0.45) = 0.036$
$W_5$	1	0	1	no	0.0
$W_6$	1	1	0	no	0.0
$W_7$	0	1	0	no	0.0
$W_8$	1	0	0	no	0.0

#SAT : 4

WMC: 0.375

$$p(\mathbf{P}_3) = 0.365 / (1 - 0.365)$$

# A Probabilistic Approach: Weighted Model Counting

- **Querying:** generalizes to any propositional formula  $\alpha$  and reducible to the problem of Weighted Model Counting (WMC) (Chavira and Darwiche, 2008)

$$p_{\text{query}}(\alpha) = \underbrace{\sum_{W \text{ s.t. } W \models \alpha} p(W)}_{\text{WMC}}$$

constraints

$$\alpha = P_1 \rightarrow P_2 \wedge P_2 \rightarrow P_3$$

world $W$	$P_1$	$P_2$	$P_3$	model?	$p(W)$
$W_1$	0	0	0	yes	$(0.1 \times 0.2 \times 0.55) = 0.01$
$W_2$	1	1	1	yes	$(0.9 \times 0.8 \times 0.45) = 0.32$
$W_3$	0	0	1	yes	$(0.1 \times 0.2 \times 0.45) = 0.009$
$W_4$	0	1	1	yes	$(0.1 \times 0.8 \times 0.45) = 0.036$
$W_5$	1	0	1	yes	$(0.9 \times 0.2 \times 0.45) = 0.081$
$W_6$	1	1	0	yes	$(0.9 \times 0.8 \times 0.55) = 0.39$
$W_7$	0	1	0	yes	$(0.1 \times 0.8 \times 0.55) = 0.04$
$W_8$	1	0	0	yes	$(0.9 \times 0.2 \times 0.55) = 0.08$

$$p(\alpha) = 0.375$$



# A Probabilistic Approach: Weighted Model Counting

- **Querying:** generalizes to any propositional formula  $\alpha$  and reducible to the problem of Weighted Model Counting (WMC) (Chavira and Darwiche, 2008)

$$p_{\text{query}}(\alpha) = \underbrace{\sum_{W \text{ s.t. } W \models \alpha} p(W)}_{\text{WMC}}$$

constraints

$$\alpha = P_1 \rightarrow P_2 \wedge P_2 \rightarrow P_3$$

world $W$	$P_1$	$P_2$	$P_3$	model?	$p(W)$
$W_1$	0	0	0	yes	$(0.1 \times 0.2 \times 0.55) = 0.01$
$W_2$	1	1	1	yes	$(0.9 \times 0.8 \times 0.45) = 0.32$
$W_3$	0	0	1	yes	$(0.1 \times 0.2 \times 0.45) = 0.009$
$W_4$	0	1	1	yes	$(0.1 \times 0.8 \times 0.45) = 0.036$
$W_5$	1	0	1	yes	$(0.9 \times 0.2 \times 0.45) = 0.081$
$W_6$	1	1	0	yes	$(0.9 \times 0.8 \times 0.55) = 0.39$
$W_7$	0	1	0	yes	$(0.1 \times 0.8 \times 0.55) = 0.04$
$W_8$	1	0	0	yes	$(0.9 \times 0.2 \times 0.55) = 0.08$

$$p(\alpha) = 0.375$$


Other: **MARG** or **SUCC** inference, *query probability*), (WMC),  
**MPE/MAP**, most likely world (**MaxSAT**) (De Raedt and Kimmig, 2015).

# A Probabilistic Approach: Weighted Model Counting

- **Querying:** generalizes to any propositional formula  $\alpha$  and reducible to the problem of Weighted Model Counting (WMC) (Raedt et al., 2016)

$$p_{query}(\alpha) = \underbrace{\sum_{W \text{ s.t. } W \models \alpha} p(W)}_{\text{WMC}}$$

**Efficient** marginal computations through knowledge compilation (Darwiche and Marquis, 2002), many open-source compilers and tools.




# A Probabilistic Approach: Weighted Model Counting

- **Querying:** generalizes to any propositional formula  $\alpha$  and reducible to the problem of Weighted Model Counting (WMC) (Raedt et al., 2016)

$$p_{\text{query}}(\alpha) = \underbrace{\sum_{W \text{ s.t. } W \models \alpha} p(W)}_{\text{WMC}}$$

**Efficient** marginal computations through knowledge compilation (Darwiche and Marquis, 2002), many open-source compilers and tools.



```
from pysdd.sdd import SddManager, Vtree, WmcManager # pip install PySDD
vtree = Vtree(var_count=4, var_order=[2,1,4,3], vtree_type="balanced")
sdd = SddManager.from_vtree(vtree); a, b, c, d = sdd.vars

alpha = (a & b) | (b & c) | (c & d)
wmc = alpha.(log_mode=False); wmc.set_literal_weight(a, 0.5)
print(f"Weighted Model Count: {wmc.propagate()}")
```

# Logic as Loss Function: probabilistic variant


- Provides an alternative to fuzzy semantics, compute *marginal probabilities* of formulas  $\alpha$  over propositions  $\mathbf{P}_1, \dots, \mathbf{P}_n$

$$\mathbf{Loss}_\alpha \propto -\log \sum_{W \text{ s.t. } W \models \alpha} \underbrace{\prod_{\mathbf{P} \in W^1} p_{model}(\mathbf{P}) \cdot \prod_{\mathbf{P} \in W^0} 1 - p_{model}(\mathbf{P})}_{\text{parameterized by our model}}$$

# Logic as Loss Function: probabilistic variant

- Provides an alternative to fuzzy semantics, compute *marginal probabilities* of formulas  $\alpha$  over propositions  $\mathbf{P}_1, \dots, \mathbf{P}_n$

$$\mathbf{Loss}_\alpha \propto -\log \sum_{W \text{ s.t. } W \models \alpha} \underbrace{\prod_{\mathbf{P} \in W^1} p_{model}(\mathbf{P}) \cdot \prod_{\mathbf{P} \in W^0} 1 - p_{model}(\mathbf{P})}_{\text{parameterized by our model}}$$



# Logic as Loss Function: probabilistic variant

- Provides an alternative to fuzzy semantics, compute *marginal probabilities* of formulas  $\alpha$  over propositions  $\mathbf{P}_1, \dots, \mathbf{P}_n$

$$\mathbf{Loss}_\alpha \propto -\log \sum_{W \text{ s.t. } W \models \alpha} \underbrace{\prod_{\mathbf{P} \in W^1} p_{\text{model}}(\mathbf{P}) \cdot \prod_{\mathbf{P} \in W^0} 1 - p_{\text{model}}(\mathbf{P})}_{\text{parameterized by our model}}$$

**semantic loss function** ([Xu et al., 2018](#))

The semantic loss is proportional to a negative logarithm of the probability of generating a state that satisfies the constraint when sampling values according to  $p$ . Hence, it is the self-information (or ‘surprise’) of obtaining an assignment that satisfies the constraint...

# Logic as Loss Function: probabilistic variant

- Provides an alternative to fuzzy semantics, compute *marginal probabilities* of formulas  $\alpha$  over propositions  $\mathbf{P}_1, \dots, \mathbf{P}_n$

$$\mathbf{Loss}_\alpha \propto -\log \sum_{W \text{ s.t. } W \models \alpha} \underbrace{\prod_{\mathbf{P} \in W^1} p_{\text{model}}(\mathbf{P}) \cdot \prod_{\mathbf{P} \in W^0} 1 - p_{\text{model}}(\mathbf{P})}_{\text{parameterized by our model}}$$

## semantic loss function (Xu et al., 2018)

The semantic loss is proportional to a negative logarithm of the probability of generating a state that satisfies the constraint when sampling values according to  $p$ . Hence, it is the self-information (or ‘surprise’) of obtaining an assignment that satisfies the constraint...

As before, often used as *undirected* model (alternative (Manhaeve et al., 2018)):

$$\mathbf{Loss} = \underbrace{\mathbf{Loss}_{\text{prediction}}}_{\text{ordinary loss}} + \underbrace{\lambda \mathbf{Loss}_\alpha}_{\text{logical constraints}}$$

# The NLI Example Again

x1: A man with a hat is riding his bicycle down the street.

x2: A person is moving with the help of their legs.

Prediction: Entailment

x2: A person is moving with the help of their legs.

x3: A person is moving

Prediction: Entailment

$$\underbrace{0.95 :: \mathbf{E}(x_1, x_2) \wedge 0.85 :: \mathbf{E}(x_2, x_3) \wedge 0.75 :: \mathbf{C}(x_1, x_3) \equiv \neg \mathbf{E}(x_1, x_3)}_{\text{(weighted) model predictions}} \longrightarrow \begin{array}{l} \alpha_{\text{constraints}} \text{ (before grounding)} \\ \forall x, y, z (\mathbf{E}(x, y) \wedge \mathbf{E}(y, z) \rightarrow \mathbf{E}(x, z)) \\ \forall x, y, z (\mathbf{E}(x, y) \wedge \mathbf{C}(y, z) \rightarrow \mathbf{C}(x, z)) \\ \forall x, y (\mathbf{C}(x, y) \rightarrow \mathbf{C}(y, x)) \end{array}$$

world $W$	$\mathbf{E}(x_1, x_2)$	$\mathbf{E}(x_2, x_3)$	$\mathbf{E}(x_1, x_3)$	$W \models \alpha$	$p(W)$
$W_1$	0	0	0	yes	$(0.05 \times 0.15 \times 0.75) = 0.005$
$W_2$	1	1	1	yes	$(0.95 \times 0.85 \times 0.25) = 0.201$
$W_3$	0	0	1	yes	$(0.05 \times 0.15 \times 0.25) = 0.001$
$W_4$	0	1	1	yes	$(0.05 \times 0.85 \times 0.25) = 0.010$
$W_5$	1	0	1	no	$(0.95 \times 0.15 \times 0.25) = 0.03$
$W_6$	1	1	0	no	$(0.95 \times 0.85 \times 0.75) = 0.60$
$W_7$	0	1	0	yes	$(0.05 \times 0.85 \times 0.75) = 0.03$
$W_8$	1	0	0	yes	$(0.95 \times 0.15 \times 0.75) = 0.10$

$$p(\alpha_{\text{constraints}}) \approx 0.34$$



# The NLI Example Again

x1: A man with a hat is riding his bicycle down the street.

x2: A person is moving with the help of their legs.

Prediction: Entailment

x2: A person is moving with the help of their legs.

x3: A person is moving

Prediction: Entailment

$$0.95 :: E(x_1, x_2) \wedge 0.85 :: E(x_2, x_3) \wedge 0.75 :: C(x_1, x_3) \equiv \neg E(x_1, x_3)$$

(weighted) model predictions,  $\alpha_{\text{pred}} = E(x_1, x_2) \wedge E(x_2, x_3) \wedge C(x_1, x_3)$

$\alpha_{\text{constraints}}$  (before grounding)

$$\forall x, y, z (E(x, y) \wedge E(y, z) \rightarrow E(x, z))$$

$$\forall x, y, z (E(x, y) \wedge C(y, z) \rightarrow C(x, z))$$

$$\forall x, y (C(x, y) \rightarrow C(y, x))$$

world $W$	$E(x_1, x_2)$	$E(x_2, x_3)$	$E(x_1, x_3)$	$W \models \alpha$	$p(W)$
$W_1$	0	0	0	no	$(0.05 \times 0.15 \times 0.75) = 0.005$
$W_2$	1	1	1	no	$(0.95 \times 0.85 \times 0.25) = 0.201$
$W_3$	0	0	1	no	$(0.05 \times 0.15 \times 0.25) = 0.001$
$W_4$	0	1	1	no	$(0.05 \times 0.85 \times 0.25) = 0.010$
$W_5$	1	0	1	no	$(0.95 \times 0.15 \times 0.25) = 0.03$
$W_6$	1	1	0	no	violates constraints
$W_7$	0	1	0	no	$(0.05 \times 0.85 \times 0.75) = 0.03$
$W_8$	1	0	0	no	$(0.95 \times 0.15 \times 0.75) = 0.10$

$$p(\alpha_{\text{constraints}} \wedge \alpha_{\text{pred}}) = 0.0$$

# The NLI Example Again

x1: A man with a hat is riding his bicycle down the street.

x2: A person is moving with the help of their legs.

Prediction: Entailment

x2: A person is moving with the help of their legs.

x3: A person is moving

Prediction: Entailment

$$0.95 :: \mathbf{E}(x_1, x_2) \wedge 0.85 :: \mathbf{E}(x_2, x_3) \wedge 0.75 :: \mathbf{C}(x_1, x_3) \equiv \neg \mathbf{E}(x_1, x_3)$$

(weighted) model predictions,  $\alpha_{\text{pred}} = \mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$

$$\begin{aligned} &\forall x, y, z (\mathbf{E}(x, y) \wedge \mathbf{E}(y, z) \rightarrow \mathbf{E}(x, z)) \\ &\forall x, y, z (\mathbf{E}(x, y) \wedge \mathbf{C}(y, z) \rightarrow \mathbf{C}(x, z)) \\ &\forall x, y (\mathbf{C}(x, y) \rightarrow \mathbf{C}(y, x)) \end{aligned}$$

world $W$	$\mathbf{E}(x_1, x_2)$	$\mathbf{E}(x_2, x_3)$	$\mathbf{E}(x_1, x_3)$	$W \models \alpha$	$p(W)$
$W_1$	0	0	0	no	$(0.05 \times 0.15 \times 0.75) = 0.005$
$W_2$	1	1	1	no	$(0.95 \times 0.85 \times 0.25) = 0.201$
$W_3$	0	0	1	no	$(0.05 \times 0.15 \times 0.25) = 0.001$
$W_4$	0	1	1	no	$(0.05 \times 0.85 \times 0.25) = 0.010$
$W_5$	1	0	1	no	$(0.95 \times 0.15 \times 0.25) = 0.03$
$W_6$	1	1	0	yes	$(0.95 \times 0.85 \times 0.75) = 0.60$
$W_7$	0	1	0	no	$(0.05 \times 0.85 \times 0.75) = 0.03$
$W_8$	1	0	0	no	$(0.95 \times 0.15 \times 0.75) = 0.10$

$$p(\alpha_{\text{pred}}) = 0.60$$

# The NLI Example Again

x1: A man with a hat is riding his bicycle down the street.

x2: A person is moving with the help of their legs.

Prediction: Entailment

x2: A person is moving with the help of their legs.

x3: A person is moving

Prediction: Entailment

$$0.95 :: \mathbf{E}(x_1, x_2) \wedge 0.85 :: \mathbf{E}(x_2, x_3) \wedge 0.75 :: \mathbf{C}(x_1, x_3) \equiv \neg \mathbf{E}(x_1, x_3)$$

(weighted) model predictions,  $\alpha_{\text{pred}} = \mathbf{E}(x_1, x_2) \wedge \mathbf{E}(x_2, x_3) \wedge \mathbf{C}(x_1, x_3)$

$$\begin{aligned} \forall x, y, z (\mathbf{E}(x, y) \wedge \mathbf{E}(y, z) \rightarrow \mathbf{E}(x, z)) \\ \forall x, y, z (\mathbf{E}(x, y) \wedge \mathbf{C}(y, z) \rightarrow \mathbf{C}(x, z)) \\ \forall x, y (\mathbf{C}(x, y) \rightarrow \mathbf{C}(y, x)) \end{aligned}$$

world $W$	$\mathbf{E}(x_1, x_2)$	$\mathbf{E}(x_2, x_3)$	$\mathbf{E}(x_1, x_3)$	$W \models \alpha$	$p(W)$
$W_1$	0	0	0	no	$(0.05 \times 0.15 \times 0.75) = 0.005$
$W_2$	1	1	1	no	$(0.95 \times 0.85 \times 0.25) = 0.201$
$W_3$	0	0	1	no	$(0.05 \times 0.15 \times 0.25) = 0.001$
$W_4$	0	1	1	no	$(0.05 \times 0.85 \times 0.25) = 0.010$
$W_5$	1	0	1	no	$(0.95 \times 0.15 \times 0.25) = 0.03$
$W_6$	1	1	0	yes	$(0.95 \times 0.85 \times 0.75) = 0.60$
$W_7$	0	1	0	no	$(0.05 \times 0.85 \times 0.75) = 0.03$
$W_8$	1	0	0	no	$(0.95 \times 0.15 \times 0.75) = 0.10$

$$p(\alpha_{\text{pred}}) = 0.60$$

Ordinary loss is again special case:  $-\log p(\alpha_{\text{pred}}) = \mathbf{Loss}_{\text{ordinary loss}}$

# Conclusion

- ▶ **Neural Symbolic modeling**, focusing on the 'logic as loss function' and 'logic in weights' approaches, **undirected models**

**conceptual tools connecting logic and DL**: thinking of model predictions as symbolic objects, annotations as logical specifications

**technical tools**: fuzzy and soft logic relaxations, model counting (probabilistic approach), translating logic to loss functions.

Still a niche area in NLP, many exciting topics to explore.

# Credits and Additional Reading

- ▶ Many ideas and examples taken from the following (beyond what's cited):

Guy Van den Broeck et al. tutorial:

<https://web.cs.ucla.edu/~guyvdb/talks/IJCAI16-tutorial/>, (Fierens et al., 2015; Raedt et al., 2016; Manhaeve et al., 2021),

**Additional resources:** *Problog*: <https://dtai.cs.kuleuven.be/problog/>,  
*PySDD*: <https://github.com/wannesm/PySDD>, *Pylon*:  
<https://pylon-lib.github.io/>

Thank you.

# References I

- Asai, A. and Hajishirzi, H. (2020). Logic-guided data augmentation and regularization for consistent question answering. *arXiv preprint arXiv:2004.10157*.
- Chavira, M. and Darwiche, A. (2008). On probabilistic inference by weighted model counting. *Artificial Intelligence*, 172(6-7):772–799.
- Darwiche, A. and Marquis, P. (2002). A knowledge compilation map. *Journal of Artificial Intelligence Research*, 17:229–264.
- De Raedt, L. and Kimmig, A. (2015). Probabilistic (logic) programming concepts. *Machine Learning*, 100(1):5–47.
- Fierens, D., Van den Broeck, G., Renkens, J., Shterionov, D., Gutmann, B., Thon, I., Janssens, G., and De Raedt, L. (2015). Inference and learning in probabilistic logic programs using weighted boolean formulas. *Theory and Practice of Logic Programming*, 15(3):358–401.
- Grespan, M. M., Gupta, A., and Srikumar, V. (2021). Evaluating relaxations of logic for neural networks: A comprehensive study. *arXiv preprint arXiv:2107.13646*.
- Li, T., Gupta, V., Mehta, M., and Srikumar, V. (2019). A logic-driven framework for consistency of neural models. *arXiv preprint arXiv:1909.00126*.
- Manhaeve, R., Dumancic, S., Kimmig, A., Demeester, T., and De Raedt, L. (2018). Deepproblog: Neural probabilistic logic programming. *Advances in Neural Information Processing Systems*, 31.

# References II

- Manhaeve, R., Dumančić, S., Kimmig, A., Demeester, T., and De Raedt, L. (2021). Neural probabilistic logic programming in deepprolog. *Artificial Intelligence*, 298:103504.
- Marra, G., Dumančić, S., Manhaeve, R., and De Raedt, L. (2021). From statistical relational to neural symbolic artificial intelligence: a survey. *arXiv preprint arXiv:2108.11451*.
- Minervini, P. and Riedel, S. (2018). Adversarially regularising neural nli models to integrate logical background knowledge. *arXiv preprint arXiv:1808.08609*.
- Raedt, L. D., Kersting, K., Natarajan, S., and Poole, D. (2016). Statistical relational artificial intelligence: Logic, probability, and computation. *Synthesis lectures on artificial intelligence and machine learning*, 10(2):1–189.
- Rocktäschel, T., Singh, S., and Riedel, S. (2015). Injecting logical background knowledge into embeddings for relation extraction. In *Proceedings of the 2015 conference of the north American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 1119–1129.
- Xu, J., Zhang, Z., Friedman, T., Liang, Y., and Broeck, G. (2018). A semantic loss function for deep learning with symbolic knowledge. In *International conference on machine learning*, pages 5502–5511. PMLR.